

Kubernetes

Minions: This is an individual node used in Kubernetes

Combination of these minions is called as Kubernetes cluster

Master is the main machine which triggers the container orchestration

It distributes the work load to the Slaves

Slaves are the nodes that accept the work load from the master

and handle activities load balancing, autoscaling, high availability etc

Kubernetes uses various types of Objects

1 Pod: This is a layer of abstraction on top of a container. This is the smallest object that Kubernetes can work on. In the Pod we have a container.

The advantage of using a Pod is that `kubectl` commands will work on the Pod and the

Pod communicates these instructions to the container. In this way we can use the same `kubectl` irrespective of which technology containers are in the Pod.

2 Service: This is used for port mapping and network load balancing

3 Namespace: This is used for creating partitions in the cluster. Pods running in a namespace cannot communicate with other pods running in other namespaces

4 Secrets: This is used for passing encrypted data to the Pods

5 ReplicationController: This is used for managing multiple replicas of PODs

and also performing scaling

6 ReplicaSet: This is similar to replicationcontroller but it is more advanced where features like selector can be implemented

7 Deployment: This used for performing all activities that a Replicaset can do it can also handle rolling update

8 PersistentVolume: Used to specify the section of storage that should be used for volumes

9 PersistentVolumeClaims: Used to reserve a certain amount of storage for a pod from the persistent volume.

10 Statefulsets: These are used to handle stateful application like data bases where consistency in read write operations has to be maintained.

11 HorizontalPodAutScaler: Used for auto scaling of pods depending on the load

Kubernetes Architecture

Master Componentes

Container runtime: This can be docker or anyother container technology

apiServer: Users interact with the apiServer using some client like ui,command line tool like kubelet.It is the apiServer which is the gateway to the cluster It works as a gatekeeper for authentication and it validates if a specific user is having permissions to execute a specific command.Example if we want to deploy a pod or a deployment first apiServers validates if the user is authorised to perform that action and if so it passes to the next process

ie the "Scheduler"

Scheduler: This process accepts the instructions from apiServer after validation and starts an application on a specific node or set of nodes. It estimates how much amount of h/w is required for an application and then checks which slave have the necessary h/w resources and instructs the kubelet to deploy the application

kubelet: This is the actual process that takes the orders from scheduler and deploy an application on a slave. This kubelet is present on both master and slave

controller manager: This checks if the desired state of the cluster is always maintained. If a pod dies it recreates that pod to maintain the desired state

etcd: Here the cluster state is maintained in key value pairs. It maintains info about the slaves and the h/w resources available on the slaves and also the pods running on the slaves. The scheduler and the control manager read the info from this etcd and schedule the pods and maintain the desired state

=====

===

Worker components

=====

containerrun time: Docker or some other container technology

kubelet: This process interacts with container run time and the node and it starts a pod with a container in it

kubeproxy: This will take the request from services to pod

It has the intelligence to forward a request to a nearby pod. Eg If an application pod wants to communicate with a db pod then kubeproxy will take that request to the nearby pod

Kubernetes setup can be done in 2 ways

1 Managed Kubernetes setup

On AWS (EKS : Elastic Kubernetes service)

On GCP (GKE : Google Kubernetes Engine)

On Azure (AKS : Azure Kubernetes Service)

2 Unmanaged Kubernetes setup

1 KOPS

2 Kubeadm

KOPS installations

Kubernetes on AWS using Kops

1. Launch Linux EC2 instance in AWS (Kubernetes Client)

2. Create and attach IAM role to EC2 Instance.

Kops need permissions to access

S3

EC2

VPC

Route53

Autoscaling

etc..

3. Install Kops on EC2

```
curl -LO https://github.com/kubernetes/kops/releases/download/$(curl -s  
https://api.github.com/repos/kubernetes/kops/releases/latest | grep tag_name |  
cut -d '"' -f 4)/kops-linux-amd64
```

```
chmod +x kops-linux-amd64
```

```
sudo mv kops-linux-amd64 /usr/local/bin/kops
```

4. Install kubectl

```
curl -LO https://storage.googleapis.com/kubernetes-release/release/$(curl -s
```

`https://storage.googleapis.com/kubernetes-release/release/stable.txt)/bin/linux/amd64/kubectl`

`chmod +x ./kubectl`

`sudo mv ./kubectl /usr/local/bin/kubectl`

5. Create S3 bucket in AWS

S3 bucket is used by kubernetes to persist cluster state, lets create s3 bucket using aws cli Note: Make sure you choose bucket name that is unique accross all aws accounts

`aws s3 mb s3://project.in.k8s --region us-west-2`

6. Create private hosted zone in AWS Route53

Head over to aws Route53 and create hostedzone

Choose name for example (sai.in)

Choose type as privated hosted zone for VPC

Select default vpc in the region you are setting up your cluster

Hit create

7 Configure environment variables.

Open .bashrc file

`vi ~/.bashrc`

Add following content into .bashrc, you can choose any arbitrary name for cluster and make sure buck name matches the one you created in previous step.

`export KOPS_CLUSTER_NAME=intelliqq.in`

`export KOPS_STATE_STORE=s3://intelliqq.in.k8s`

Then running command to reflect variables added to .bashrc

`source ~/.bashrc`

8. Create ssh key pair

This keypair is used for ssh into kubernetes cluster

`ssh-keygen`

9. Create a Kubernetes cluster definition.

`kops create cluster \`

```
--state=${KOPS_STATE_STORE} \  
--node-count=2 \  
--master-size=t3.medium \  
--node-size=t3.medium \  
--zones=us-west-2a \  
--name=${KOPS_CLUSTER_NAME} \  
--dns private \  
--master-count 1
```

10. Create kubernetes cluster

```
kops update cluster --yes --admin
```

Above command may take some time to create the required infrastructure resources on AWS. Execute the validate command to check its status and wait until the cluster becomes ready

```
kops validate cluster
```

For the above command, you might see validation failed error initially when you create cluster and it is expected behaviour, you have to wait for some more time and check again.

11. To connect to the master

```
ssh admin@api.javahome.in
```

Destroy the kubernetes cluster

```
kops delete cluster --yes
```

Update Nodes and Master in the cluster

We can change number of nodes and number of masters using following commands

```
kops edit ig nodes change minSize and maxSize to 0
```

```
kops get ig- to get master node name
```

```
kops edit ig - change min and max size to 0
```

```
kops update cluster --yes
```

=====

=

